

Applied Economic Forecasting

2. Exploring & Visualizing Time series

Spring 2020

- 1 Time series in R
- 2 Time plots
- 3 Seasonal plots
- 4 Seasonal or cyclic?
- 5 Scatterplots
- 6 Lag plots and autocorrelation
- 7 White noise

Section 1

Time series in R

ts objects and ts function

A time series is stored in a `ts` object in R:

- a list of numbers
- information about times those numbers were recorded.

Example

Year	Obs.
2012	123
2013	39
2014	78
2015	52
2016	110

```
y <- ts(c(123,39,78,52,110), start=2012)
```

ts objects and ts function

For observations that are more frequent than once per year, we need to add a **frequency** argument.

E.g., monthly data stored as a numerical vector **z**:

```
y <- ts(z, frequency=12, start=c(2003, 1))
```

ts objects and ts function

`ts(data, frequency, start)`

Type of data	frequency	start example
Annual	1	1995
Quarterly	4	c(1995,2)
Monthly	12	c(1995,9)
Daily	7 or 365.25	or c(1995,234)
Weekly	52.18	c(1995,23)
Hourly	24 or 168 or 8,766	1
Half-hourly	48 or 336 or 17,532	1

Let's Practice!!!

- ① Set a seed as 10^3
- ② Generate a random normal variable (\mathbf{x}) with 200 observations, mean=75, and sd= 5
- ③ Declare \mathbf{x} as a *quarterly* ts object ending December 2018 ($\mathbf{x.ts}$)
- ④ Now repeat step 3 with weekly, monthly and annual frequencies (*How about using a loop?*)

Australian GDP

```
ausgdp <- ts(x, frequency=4, start=c(1971,3))
```


- Class: “ts”
- Print and plotting methods available.

```
head(ausgdp, 30)
```

##	Qtr1	Qtr2	Qtr3	Qtr4
## 1971			4612	4651
## 1972	4645	4615	4645	4722
## 1973	4780	4830	4887	4933
## 1974	4921	4875	4867	4905
## 1975	4938	4934	4942	4979
## 1976	5028	5079	5112	5127
## 1977	5130	5101	5072	5069
## 1978	5100	5166	5244	5312

Australian GDP

```
autoplot(ausgdp)
```



Residential electricity sales

```
elecsales
```

```
## Time Series:
```

```
## Start = 1989
```

```
## End = 2008
```

```
## Frequency = 1
```

```
## [1] 2354.34 2379.71 2318.52 2468.99 2386.09 2569.47 2
```

```
## [10] 3000.70 3108.10 3357.50 3075.70 3180.60 3221.60 3
```

```
## [19] 3637.89 3655.00
```

```
> library(fpp2)
```

This loads:


- some data for use in examples and exercises
- **forecast** package (for forecasting functions)
- **ggplot2** package (for graphics functions)
- **fma** package (for lots of time series data)
- **expsmooth** package (for more time series data)

Section 2

Time plots


Time plots

```
autoplot(melsyd[, "Economy.Class"])
```



Time plots

```
autoplot(a10) + ylab("$ million") + xlab("Year") +  
  ggtitle("Antidiabetic drug sales")
```




Let's Practice

- ① Using the data from our earlier exercise: Plot `x.ts` at the monthly, quarterly and yearly frequencies.
- ② Create plots of the following time series: `dole`, `bricksq`, `lynx`, `goog`
 - Use `help()` to find out about the data in each series.
 - For each of the plots, be sure to modify the axis labels and title.


Are time plots best?

```
autoplot(elecdaily[, "Temperature"]) + labs(x = "Week",  
      y = "Max temperature")
```



Are time plots best?

```
qplot(time(elecdaily), elecdaily[, "Temperature"]) +  
  labs(x = "Week", y = "Max temperature")
```




Section 3

Seasonal plots

Seasonal plots

```
ggseasonplot(a10, year.labels = TRUE, year.labels.left = TRUE) +  
  ylab("$ million") + ggtitle("Seasonal plot: antidiabetic drug sales")
```

Seasonal plot: antidiabetic drug sales




- Data plotted against the individual “seasons” in which the data were observed. (In this case a “season” is a month.)
- Something like a time plot except that the data from each season are overlapped.
- Enables the underlying seasonal pattern to be seen more clearly, and also allows any substantial departures from the seasonal pattern to be easily identified.
- In R: `ggseasonplot()`

Seasonal polar plots

```
ggseasonplot(a10, polar = TRUE) + ylab("$ million")
```


Seasonal plot: a10



Seasonal subseries plots

```
ggsubseriesplot(a10) + ylab("$ million") + ggtitle("Subseries plot: ")
```

Subseries plot: antidiabetic drug sales




Seasonal subseries plots

- Data for each season collected together in time plot as separate time series.
- Enables the underlying seasonal pattern to be seen clearly, and changes in seasonality over time to be visualized.
- In R: `ggsubseriesplot()`

Quarterly Australian Beer Production

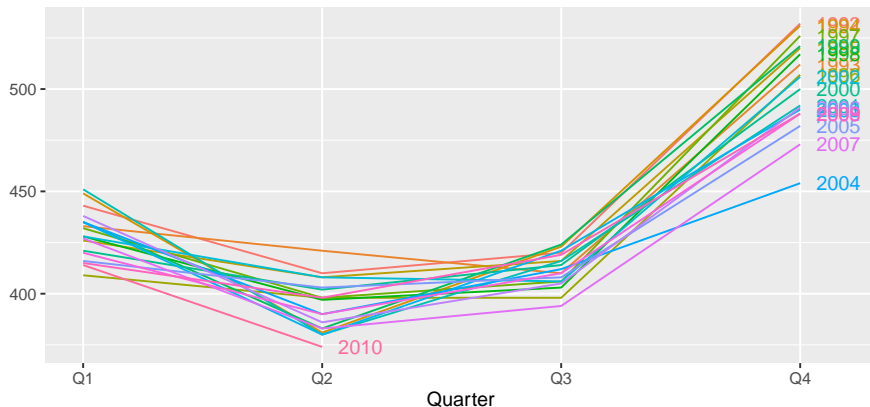
```
beer <- window(ausbeer, start = 1992)  
autoplot(beer)
```



Quarterly Australian Beer Production


```
ggseasonplot(beer, year.labels = TRUE)
```

Seasonal plot: beer



Quarterly Australian Beer Production

```
ggsubseriesplot(beer)
```



Let's Practice!!!

The `arrivals` data set comprises quarterly international arrivals (in thousands) to Australia from Japan, New Zealand, UK and the US.

- Use `autoplot()` and `ggseasonplot()` to compare the differences between the arrivals from these four countries.
- Can you identify any unusual observations?

Section 4

Seasonal or cyclic?

Time series patterns

Trend pattern exists when there is a long-term increase or decrease in the data.

Seasonal pattern exists when a series is influenced by seasonal factors (e.g., the quarter of the year, the month, or day of the week).

Cyclic pattern exists when data exhibit rises and falls that are *not of fixed period* (duration usually of at least 2 years).

Time series components


Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

Time series patterns


```
autoplot(window(elec, start = 1980)) + ggtitle("Australian electricity  
  labs(x = "Year", y = "GWh")
```

Australian electricity production



Time series patterns


```
autoplot(bricksq) + ggtitle("Australian clay brick production") +  
  xlab("Year") + ylab("million units")
```



Time series patterns


```
autoplot(hsales) + ggtitle("Sales of new one-family houses, USA") +  
  labs(x = "Year", y = "Total sales")
```

Sales of new one-family houses, USA




Time series patterns

```
autoplot(ustreas) + ggtitle("US Treasury Bill Contracts") +  
  labs(x = "Day", y = "price")
```



Time series patterns

```
autoplot(lynx) + ggtitle("Annual Canadian Lynx Trappings") +  
  xlab("Year") + ylab("Number trapped")
```



Seasonal or cyclic?

Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

The timing of peaks and troughs is predictable with seasonal data, but unpredictable in the long term with cyclic data.


Section 5

Scatterplots

Scatterplots: Relationship *Between* Time Series

Scatterplots are commonly used tools for visualizing the relationship between variables. Take the text example of Electricity Demand and Temperature in Victoria, Australia:

```
qplot(x = Temperature, y = Demand, data = as.data.frame(elecddemand)) +  
  labs(y = "Demand", x = "Temperature (C)")
```



Scatterplots: Relationship *Between* Time Series

We notice that higher temperatures are associated with higher demand for electricity.

What if we would like to get a measure of the strength of this relationship?

- For this, we will use the **correlation coefficient**:

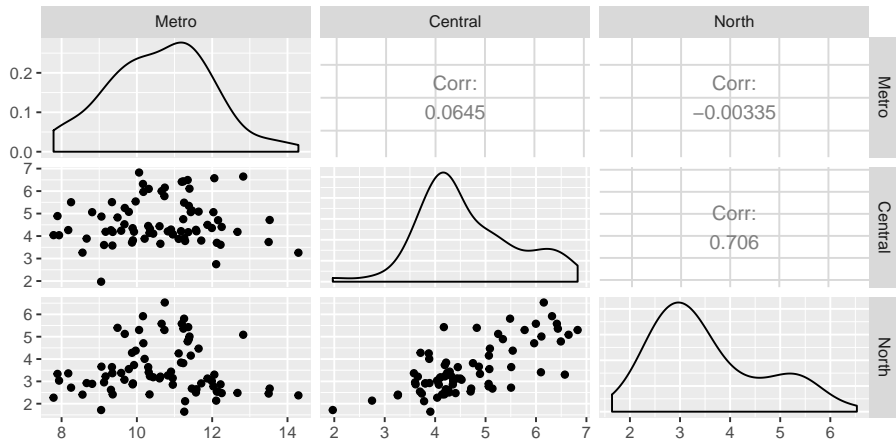
$$\rho_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sigma_x^2} \sqrt{\sigma_y^2}}, \quad -1 \leq \rho_{xy} \leq 1$$

- Negative values indicate a negative *linear* relationship between x and y .
- Positive values indicate a positive *linear* relationship between x and y .

Scatterplot Matrices: Relationship *Between* Time Series

We can visualize the correlation between multiple series using a scatterplot matrix. In R, we will use the `ggpairs()` command from the `GGally`

```
# install.packages('GGally') #unmute when running  
# for the first time  
GGally::ggpairs(as.data.frame(visnights[, 6:8]), columnLabels = c("Metro",  
  "Central", "North"))
```

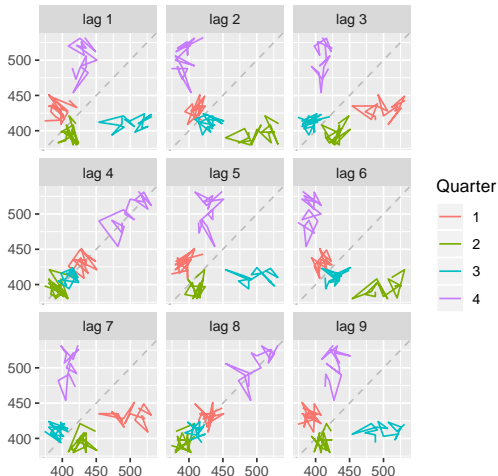


Section 6

Lag plots and autocorrelation

Example: Beer production

```
beer <- window(ausbeer, start = 1992)
gglagplot(beer)
```



Lagged scatterplots

- Each graph shows y_t plotted against y_{t-k} for different values of k .
- The autocorrelations are the correlations associated with these scatterplots.

Covariance and **correlation**: measure extent of **linear relationship** between two variables (y and X).

Autocovariance and **autocorrelation**: measure linear relationship between **lagged values** of a time series y .

We measure the relationship between:

- y_t and y_{t-1}
- y_t and y_{t-2}
- y_t and y_{t-3}
- etc.

Autocorrelation

We denote the sample autocovariance at lag k by γ_k and the sample autocorrelation at lag k by ρ_k . Then define

$$\gamma_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

and

$$\rho_k = \frac{\gamma_k}{\gamma_0}$$

- It is easy to see that γ_0 is the variance of y . Let $k = 0$ then:

$$\gamma_0 = \frac{1}{T} \sum_{t=1}^T (y_t - \bar{y})(y_t - \bar{y}) = \frac{1}{T} \sum_{t=1}^T (y_t - \bar{y})^2$$


- ρ_1 indicates how successive values of y relate to each other
- ρ_2 indicates how y values two periods apart relate to each other
- ρ_k is *almost* the same as the sample correlation between y_t and y_{t-k} .

Autocorrelation

Results for first 9 lags for beer data:


ρ_1	ρ_2	ρ_3	ρ_4	ρ_5	ρ_6	ρ_7	ρ_8	ρ_9
-0.102	-0.657	-0.060	0.869	-0.089	-0.635	-0.054	0.832	-0.108

```
ggAcf(beer)
```



- ρ_4 higher than for the other lags. This is due to **the seasonal pattern in the data**: the peaks tend to be **4 quarters** apart and the troughs tend to be **2 quarters** apart.
- ρ_2 is more negative than for the other lags because troughs tend to be 2 quarters behind peaks.
- Together, the autocorrelations at lags 1, 2, ..., make up the *autocorrelation function* or ACF.
- The plot is known as a **correlogram**

```
ggAcf(beer, col = "red")
```




Trend and seasonality in ACF plots

- When data have a trend, the autocorrelations for small lags tend to be large and positive.
- When data are seasonal, the autocorrelations will be larger at the seasonal lags (i.e., at multiples of the seasonal frequency)
- When data are trended and seasonal, you see a combination of these effects.


Aus monthly electricity production

```
elec2 <- window(elec, start = 1980)  
autoplot(elec2)
```



Aus monthly electricity production

```
ggAcf(elec2, lag.max = 48)
```



Aus monthly electricity production


Time plot shows clear trend and seasonality.

The same features are reflected in the ACF.

- The slowly decaying ACF indicates trend.
- The ACF peaks at lags 12, 24, 36, ..., indicate seasonality of length 12.


Google stock price

```
autoplot(goog, col = "blue")
```



Google stock price

```
ggAcf(goog, lag.max = 100)
```



Let's Practice!!!

We have introduced the following graphics functions:


- `gglagplot`
- `ggAcf`

Explore the following time series using these functions. Can you spot any seasonality, cyclicity and trend? What do you learn about the series?


- `hsales`
- `usdeaths`
- `bricksq`
- `sunspotarea`
- `gasoline`

Which is which?


1. Daily temperature of cow




2. Monthly accidental deaths




3. Monthly air passengers




4. Annual mink trappings




A




B



C



D




Section 7


White noise

Example: White noise

```
set.seed(14567)
wn <- ts(rnorm(36))
autoplot(wn) + ggtitle("White Noise")
```



Example: White noise



ρ_1	ρ_2	ρ_3	ρ_4	ρ_5	ρ_6	ρ_7
-0.099	0.232	0.063	0.006	-0.103	0.158	-0.291

We expect each autocorrelation to be close to zero.


Sampling distribution of ρ_k for white noise data is asymptotically $N(0, \frac{1}{T})$.

- 95% of all ρ_k for white noise must lie within $\pm \frac{1.96}{\sqrt{T}}$.
- If this is not the case, the series is probably not WN.
- Common to plot lines at $\pm 1.96/\sqrt{T}$ when plotting ACF. These are the *critical values*.

Autocorrelation

Example:

```
ggAcf(wn)
```



$T = 36$ and so critical values at $\pm \frac{1.96}{\sqrt{36}} = \pm 0.327$.

All autocorrelation coefficients lie within these limits. Data cannot be distinguished from white noise.


Example: Pigs slaughtered

```
pigs2 <- window(pigs, start = 1990)
autoplot(pigs2) + xlab("Year") + ylab("thousands") +
  ggtitle("Number of pigs slaughtered in Victoria")
```



Example: Pigs slaughtered

```
ggAcf(pigs2)
```



Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 1990 through August 1995. (Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows some significant autocorrelation at lags 1, 2, and 3.
- ρ_{12} relatively large although not significant. This may indicate some slight seasonality.

These show the series is **not a white noise series**.

Let's Practice!!!

- ① Compute the daily changes in the Google stock price (`ddgoog`).
- ② Does `ddgoog` look like white noise?
 - Present a time plot of `ddgoog`
 - Present the ACF plot


Let's Practice (Solution)

Manually:


```
# initialize the ddgoog variable to be a object  
# with NAs  
ddgoog <- NA  
  
# We can use the loop to do the changes  
# calculations Remember with differencing we lose  
# the first obs so we start the counter at 2  
for (i in 2:length(goog)) {  
  ddgoog[i] <- goog[i] - goog[i - 1] # Store results in ddgoog at position [i]  
}  
# Declare as a ts and drop first obs since it's NA  
ddgoog <- ts(ddgoog[-1], start = c(2))  
round(head(ddgoog), 3)
```

```
## Time Series:  
## Start = 2  
## End = 7  
## Frequency = 1  
## [1] -0.318  4.794  0.705  2.479  7.606  8.495
```

```
autoplot(ddgoog) + ggtitle("Daily Changes in Google Stock Prices") +  
  labs(y = "Daily Changes", x = "")
```



```
ggAcf(ddgoog)
```



Let's Practice (Solution)

Using the fpp2 package:

```
dgoog <- diff(goog)
round(head(dgoog), 3)
```

```
## Time Series:
```


```
## Start = 2
```

```
## End = 7
```

```
## Frequency = 1
```

```
## [1] -0.318  4.794  0.705  2.479  7.606  8.495
```

```
autoplot(dgoog) + ggtitle("Daily Chances in Google Stock Prices") +
  labs(y = "Daily Changes", x = "")
```



Let's Practice (Solution)

So, what is your conclusion?